

# YUN-CHI PANG

Santa Clara, CA | 669-274-6138 | [yunchipang@gmail.com](mailto:yunchipang@gmail.com) | [linkedin.com/in/yunchipang](https://linkedin.com/in/yunchipang) | [github.com/yunchipang](https://github.com/yunchipang)

## EDUCATION

Northeastern University, MSc in Computer Science	09/2022–05/2025
The Chinese University of Hong Kong, BBA in Business Analytics	09/2017–12/2021

## TECHNICAL SKILLS

**Languages & Backend:** Python, Java, Rust, SQL, Django, FastAPI, Spring Boot, Node.js  
**Databases & Storage:** MySQL, PostgreSQL, MongoDB, Hive, Iceberg, HDFS, ClickHouse  
**Cloud & DevOps:** AWS (EC2, S3), GCP, Docker, Kubernetes, Jenkins, GitHub Actions  
**Testing & Developer Tools:** Git, Pytest, JUnit, Postman, Swagger/OpenAPI

## OPEN-SOURCE

- Apache Kafka** 🔄 *15+ PRs* Event streaming platform used by 80%+ of Fortune 100 03/2025–
- Modernized core modules by migrating **Scala** to **Java**, improving maintainability across 200K+ deployments
  - Removed redundant data structures and calls, improving response time by 10% and cutting config-related CPU use by 15% through algorithm and **memory optimizations**
  - Enhanced log compaction validation by adding compression parameters and extending test coverage by 30%
- Apache Airflow** 🔄 *10+ PRs* Workflow orchestration platform used by 77K+ orgs 03/2025–
- Reduced enterprise attack surface by shifting CLI from DB access to **API-driven authentication**
- Apache Gravitino** 🔄 *20+ PRs* Federated metadata lake used by Uber, Pinterest, Xiaomi 04/2025–
- Built extensible **REST APIs** for **Java**, **Python**, and **CLI** clients, enabling federation across **Hadoop**-compatible systems (**HDFS**, **Hive**, **Iceberg**) and boosting query performance for 45K+ monthly users
  - Unified metadata authorization with permission pushdown for **relational DBs**, **data warehouses** and **lake catalogs**, cutting API-layer auth logic by 60%
- Apache Hudi** 🔄 *5+ PRs* Data lake platform used by Uber, Amazon, ByteDance 07/2025–
- Boosted scalability by adding APIs for file slice splitting and incremental change processing, enabling parallel reads across 1000+ partitions and cutting query latency by 40% on large datasets
- OpenLineage** 🔄 *Open standard for data lineage adopted by Databricks, Astronomer* 07/2025–
- Ensured **Spark 4** backward compatibility via reflection-based factory patterns to handle constructor changes across 20+ connectors including Kusto and Databricks

## PROFESSIONAL EXPERIENCE

- Software Engineer Intern** – OurDate 07/2023–10/2023  
*AI-powered growth-stage social platform focused on event-driven matchmaking*
- Built scalable REST APIs (**Django**, **MySQL**) powering core matchmaking features, improving system throughput by 25% and reducing latency for 100+ daily active users
  - Led peer code reviews, mentoring junior engineers and raising code quality, cutting post-release bugs by 20%
- Data Engineer** – Dashmote 01/2022–06/2022  
*Series B SaaS data analytics company serving Fortune 500 clients across 50+ global markets*
- Built and scaled automated **AWS** data pipelines processing 10M+ records daily, reducing runtime by 88%
  - Developed a distributed scraping framework (**Python**, **Scrapy**) to ingest 1M+ restaurant menu records per day, expanding product coverage by 3x across global markets
- Ops Intern** – ShopBack 09/2020–07/2021  
*Regional e-commerce scale-up with 35M+ users across Asia-Pacific*
- Automated workflows by replacing manual tasks with **Python** tools, increasing ops efficiency by 90%
  - Eliminated 98% of support tickets by implementing Selenium-based tracking, driving 1% weekly growth