

YUN-CHI PANG

Santa Clara, CA | 669-274-6138 | yunchipang@gmail.com | linkedin.com/in/yunchipang | github.com/yunchipang

EDUCATION

Northeastern University, MSc in Computer Science	09/2022–05/2025
The Chinese University of Hong Kong, BBA in Business Analytics	09/2017–12/2021

TECHNICAL SKILLS

Languages & Backend: Python, Java, Rust, SQL, Django, FastAPI, Spring Boot, Node.js
Distributed Systems & Data Infra: Kafka, Airflow, Spark, Flink, Iceberg, HDFS, ClickHouse
Databases: MySQL, PostgreSQL, MongoDB
Cloud & DevOps: AWS (EC2, S3), GCP, Docker, Kubernetes, Jenkins, GitHub Actions
Testing & Developer Tools: Git, Pytest, JUnit, Postman, Swagger/OpenAPI

OPEN-SOURCE

Apache Kafka  15+ PRs Event streaming platform	03/2025–
<ul style="list-style-type: none">Migrated core Scala modules to Java, removing tech debt and reducing contributor onboarding time by 50%Optimized config API, reducing CPU usage by 20% and memory allocations by 60% through algorithmic tuning and lazy evaluation	
Apache Hudi  10+ PRs Data lake platform	07/2025–
<ul style="list-style-type: none">Added file-slice partitioning APIs enabling 8x faster incremental processing across distributed ETL workersIntroduced column-stats data skipping, reducing scanned data by 90% and accelerating selective queries by 10xReduced query latency by up to 100x via automatic timestamp-to-partition filter transformation, eliminating 95% of unnecessary scans	
Apache Gravitino  20+ PRs Federated metadata lake	04/2025–
<ul style="list-style-type: none">Cut authorization overhead by 60% through permission pushdown to catalogsReduced integration test boilerplate by 60% through infra refactoring and ID autogeneration	

Additional contributions across major data-infra ecosystems (Airflow, ClickHouse, OpenLineage, Sedona, OpenDAL, Ruff), focusing on reliability, performance, and developer experience.

PROFESSIONAL EXPERIENCE

Software Engineer Intern @ OurDate	07/2023–10/2023
<i>AI-powered growth-stage social platform focused on date-planning</i>	
<ul style="list-style-type: none">Boosted system throughput by 25% by developing scalable REST APIs (Django, MySQL) powering core date-planning features—reduced latency for 100+ daily active usersSuppressed post-release defects by 20% through systematic peer code reviews and mentoring junior engineers	
Data Engineer @ Dashmote	01/2022–06/2022
<i>Series B SaaS data analytics company serving Fortune 500 clients across 50+ global markets</i>	
<ul style="list-style-type: none">Built and optimized AWS pipelines processing 10M+ records/day, cutting runtime by 88%Designed scalable, multi-region scraping framework (Python, Scrapy) ingesting 1M+ menu records/day across global markets	
Ops Intern @ ShopBack	09/2020–07/2021
<i>Regional e-commerce scale-up with 35M+ users across Asia-Pacific</i>	
<ul style="list-style-type: none">Increased ops efficiency by 90% by automating manual workflows with Python toolsCut 98% of support tickets through Selenium-based tracking automation—drove 1% weekly partner growth	